

Autoencoder

February 8, 2021

```
[11]: import numpy as np
import matplotlib.pyplot as plt
import torch
import time
import sys
from collections import OrderedDict
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchsummary import summary
import torchvision
import torchvision.transforms as transforms
from torch.autograd import Variable
from typing import Union, Tuple
from torch.utils.data import Dataset, DataLoader
import cv2
```

```
[ ]: from google.colab import drive
#drive.mount('/content/drive')
```

```
[12]: train_data = torchvision.datasets.MNIST('data', train=True, download=True,
↳transform=transforms.ToTensor())
```

```
[13]: train_loader = DataLoader(train_data, batch_size=len(train_data))
batch = next(iter(train_loader))
x, y = batch
mean, std = torch.mean(x), torch.std(x)
mean, std
```

```
[13]: (tensor(0.1307), tensor(0.3081))
```

```
[14]: transform = transforms.Compose([
                                transforms.ToTensor(),
                                transforms.Normalize(mean, std)
])
```

```
train_data = torchvision.datasets.MNIST('data', train=True, download=True,
    ↪transform=transform)
test_data = torchvision.datasets.MNIST('data', train=False, download=True,
    ↪transform=transform)
```

```
[15]: batch_size = 64
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=True)
```

0.1 Exercise 1

Implement an Autoencoder that encodes the MNIST dataset to a latent dimension of size $m < 784$. Use Transposed Convolutions and/or Unpooling to solve this exercise. Train the Autoencoder and plot the reconstruction training loss. Plot 5 digits (of your choice) before and after reconstruction. Do this for two different latent dimension sizes.

```
[16]: class Autoencoder(nn.Module):
    def __init__(self):
        super(Autoencoder, self).__init__()
        self.encoder = nn.Sequential(OrderedDict([
            ('conv1', nn.Conv2d(in_channels=1,
    ↪out_channels=64, kernel_size=3, stride=2, padding=1)), #64* 14*14
            ('relu1', nn.LeakyReLU()),
            ('pool1', nn.MaxPool2d(2,2)),
    ↪
            #64*7*7
            ('conv2', nn.Conv2d(in_channels =
    ↪64, out_channels = 16, kernel_size = 3, stride = 2, padding=1)), #16*4*4
            ('relu2', nn.LeakyReLU()),
            ('pool2', nn.MaxPool2d(2,1)),
    ↪
            #16*3*3
            ('conv3', nn.Conv2d(in_channels = 16,
    ↪out_channels = 16, kernel_size = 3, stride = 1, padding = 1)),
            ('relu3', nn.LeakyReLU()))])

        self.decoder = nn.Sequential(OrderedDict([
            ('convT1', nn.
    ↪ConvTranspose2d(in_channels=16, out_channels=64, kernel_size=3, stride=2,
    ↪padding=1)), #64* 5*5
            ('relu1', nn.LeakyReLU()),
            ('convT2', nn.
    ↪ConvTranspose2d(in_channels=64, out_channels=16, kernel_size=4, stride=3,
    ↪padding=1)), #16*14*14
            ('relu2', nn.LeakyReLU()),
            ('convT3', nn.
    ↪ConvTranspose2d(in_channels=16, out_channels=1, kernel_size=2, stride=2)),
    ↪
            #1*28*28
```

```
        ('tanh', nn.Tanh()))]))
```

```
def forward(self, x):  
    x1 = self.encoder(x)  
    x2 = self.decoder(x1)  
    return x2
```

```
[17]: model = Autoencoder().cuda()
```

```
[18]: # create an optimizer object  
# Adam optimizer with learning rate 1e-3  
optimizer = optim.Adam(model.parameters(), lr=1e-3)  
  
# mean-squared error loss  
criterion = nn.MSELoss()
```

```
[19]: def train(model: nn.Module, optimizer: optim.Optimizer, data: train_loader,   
→max_epochs: int, cuda=True):  
    total_loss = []  
    for epoch in range(max_epochs):  
        for batch in data:  
            img, _ = batch  
            img = Variable(img).cuda()  
            # =====forward=====  
            output = model(img)  
            loss = criterion(output, img)  
            # =====backward=====  
            optimizer.zero_grad()  
            loss.backward()  
            optimizer.step()  
            # =====log=====  
            total_loss.append(loss.item())  
            print('epoch [{}/{}], loss:{:.4f}'.format(epoch+1, max_epochs, loss.  
→item()))  
    return total_loss
```

```
[20]: model_losses = train(model, optimizer, train_loader, max_epochs=30)
```

```
epoch [1/30], loss:0.3837  
epoch [2/30], loss:0.4051  
epoch [3/30], loss:0.3983  
epoch [4/30], loss:0.3678  
epoch [5/30], loss:0.3928  
epoch [6/30], loss:0.3539
```

```
epoch [7/30], loss:0.3494
epoch [8/30], loss:0.3258
epoch [9/30], loss:0.3786
epoch [10/30], loss:0.3493
epoch [11/30], loss:0.2783
epoch [12/30], loss:0.3571
epoch [13/30], loss:0.3716
epoch [14/30], loss:0.3445
epoch [15/30], loss:0.3438
epoch [16/30], loss:0.3499
epoch [17/30], loss:0.3412
epoch [18/30], loss:0.3336
epoch [19/30], loss:0.3920
epoch [20/30], loss:0.3491
epoch [21/30], loss:0.3336
epoch [22/30], loss:0.3285
epoch [23/30], loss:0.3638
epoch [24/30], loss:0.3422
epoch [25/30], loss:0.3752
epoch [26/30], loss:0.3460
epoch [27/30], loss:0.3623
epoch [28/30], loss:0.3198
epoch [29/30], loss:0.3369
epoch [30/30], loss:0.3363
```

```
[21]: test_examples = None

with torch.no_grad():
    for batch_features in test_loader:
        batch_features = batch_features[0]
        test_examples = batch_features.cuda()
        reconstruction = model(test_examples)
        break
```

Plot digits before and after reconstruction


```
[22]: with torch.no_grad():
        number = 10
        plt.figure(figsize=(20, 4))
        for index in range(number):
            # display original
            ax = plt.subplot(2, number, index + 1)
            plt.imshow(test_examples.cpu()[index].numpy().reshape(28, 28))
            plt.gray()
            ax.get_xaxis().set_visible(False)
            ax.get_yaxis().set_visible(False)

            # display reconstruction
```

```

ax = plt.subplot(2, number, index + 1 + number)
plt.imshow(reconstruction.cpu()[index].numpy().reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()

```



Plot the loss

```


[23]: plt.plot(model_losses)
      plt.ylabel('loss')
      plt.xlabel('epoch')
      plt.title('MSE loss with our autoencoder')

```

```

[23]: Text(0.5, 1.0, 'MSE loss with our autoencoder')

```



Another autoencoder architecture with different dimensions of latent state

```
[24]: class Autoencoder2(nn.Module):
    def __init__(self):
        super(Autoencoder2, self).__init__()
        self.encoder = nn.Sequential(OrderedDict([
            ('conv1', nn.Conv2d(in_channels=1,
            ↪out_channels=64, kernel_size=3, stride=2, padding=1)),
            ('relu1', nn.LeakyReLU()),
            ('pool1', nn.MaxPool2d(2,2)),
            ('conv2', nn.Conv2d(in_channels =
            ↪64,out_channels = 8,kernel_size = 3,stride = 3, padding=1)),
            ('relu2', nn.LeakyReLU())]))

        self.pooling = nn.MaxPool2d(2,1, return_indices=True)
        self.convT = nn.ConvTranspose2d(in_channels=8, out_channels=8,
            ↪kernel_size=3, stride=1, padding=1)
        self.unpooling = nn.MaxUnpool2d(2,2)

        self.decoder = nn.Sequential(OrderedDict([
            ('convT2', nn.ConvTranspose2d(in_channels=8,
            ↪out_channels=64, kernel_size=5, stride=3, padding=1)),
            ('relu2', nn.LeakyReLU()),
            ('convT3', nn.
            ↪ConvTranspose2d(in_channels=64, out_channels=64, kernel_size=4, stride=2,
            ↪padding = 1)),
            ('relu3', nn.LeakyReLU()),
            ('convT4', nn.
            ↪ConvTranspose2d(in_channels=64, out_channels=1, kernel_size=5, stride=1,
            ↪padding = 0)),
            ('tanh', nn.Tanh())]))

    def forward(self, x):
        x1 = self.encoder(x)
        x1_pooled, indicies = self.pooling(x1)
        #x2 = self.convT(x1_pooled)
        #x2_unpooled = self.unpooling(x2,indicies)
        x2_unpooled = self.unpooling(x1_pooled,indicies)
        reconstructed = self.decoder(x2_unpooled)
        return reconstructed
```

```
[25]: model2 = Autoencoder2().cuda()
optimizer = optim.Adam(model2.parameters(), lr=1e-3)
criterion = nn.MSELoss()
```

```
[26]: model_losses2 = train(model2, optimizer, train_loader, max_epochs=30)
```

```
epoch [1/30], loss:0.5896
epoch [2/30], loss:0.6077
epoch [3/30], loss:0.5754
epoch [4/30], loss:0.5181
epoch [5/30], loss:0.5957
epoch [6/30], loss:0.5383
epoch [7/30], loss:0.6092
epoch [8/30], loss:0.5628
epoch [9/30], loss:0.5853
epoch [10/30], loss:0.5039
epoch [11/30], loss:0.5546
epoch [12/30], loss:0.5625
epoch [13/30], loss:0.5650
epoch [14/30], loss:0.5577
epoch [15/30], loss:0.5650
epoch [16/30], loss:0.5795
epoch [17/30], loss:0.5416
epoch [18/30], loss:0.5339
epoch [19/30], loss:0.5048
epoch [20/30], loss:0.5454
epoch [21/30], loss:0.5844
epoch [22/30], loss:0.5559
epoch [23/30], loss:0.6028
epoch [24/30], loss:0.5733
epoch [25/30], loss:0.5410
epoch [26/30], loss:0.5341
epoch [27/30], loss:0.5317
epoch [28/30], loss:0.5807
epoch [29/30], loss:0.5339
epoch [30/30], loss:0.5677
```

```
[27]: test_examples = None

with torch.no_grad():
    for batch_features in test_loader:
        batch_features = batch_features[0]
        test_examples = batch_features.cuda()
        reconstruction = model2(test_examples)
        break
```


```
[28]: with torch.no_grad():
        number = 10
        plt.figure(figsize=(20, 4))
        for index in range(number):
            # display original
            ax = plt.subplot(2, number, index + 1)
            plt.imshow(test_examples.cpu()[index].numpy().reshape(28, 28))
```

```

plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)

# display reconstruction
ax = plt.subplot(2, number, index + 1 + number)
plt.imshow(reconstruction.cpu()[index].numpy().reshape(28, 28))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()

```



```

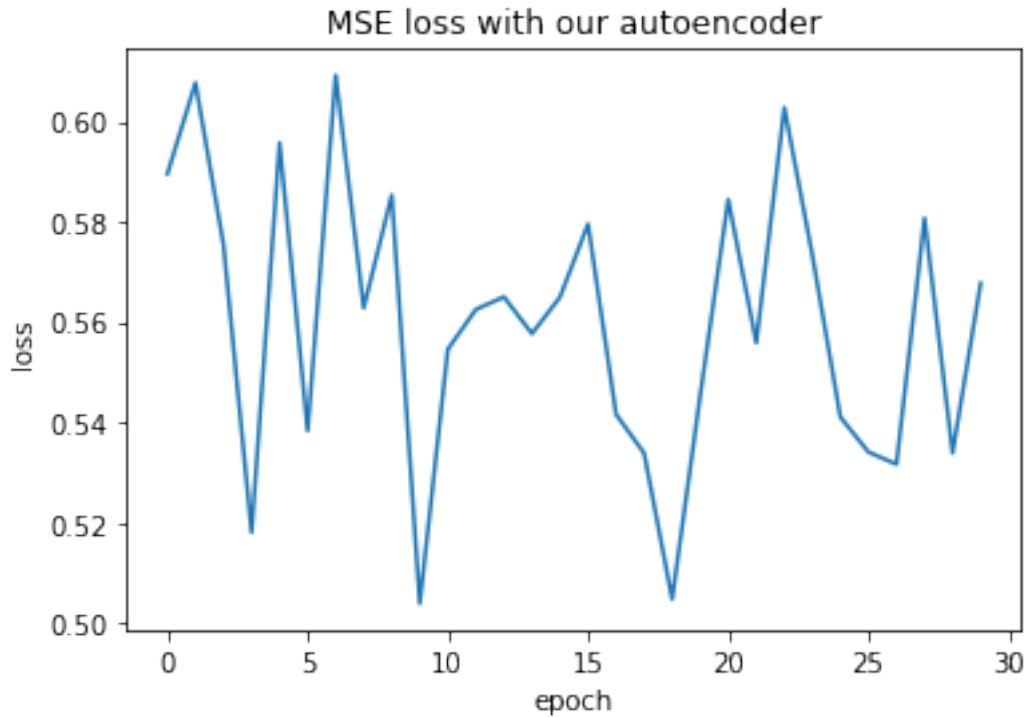
[29]: plt.plot(model_losses2)
plt.ylabel('loss')
plt.xlabel('epoch')
plt.title('MSE loss with our autoencoder')

```

```

[29]: Text(0.5, 1.0, 'MSE loss with our autoencoder')

```

0.2 Exercise 2

Now that you have built an Autoencoder, it is time to implement a Variational Autoencoder. You can use the Autoencoder you trained in the previous exercise and adapt it for this exercise. Do not forget to use the reparametrization trick for sampling from Z-space.

```
[30]: class Flatten(nn.Module):
    def forward(self, input):
        return input.view(input.size(0), -1)
class Unflatten(nn.Module):
    def __init__(self, channel, height, width):
        super(Unflatten, self).__init__()
        self.channel = channel
        self.height = height
        self.width = width

    def forward(self, input):
        return input.view(input.size(0), self.channel, self.height, self.width)

class ConvVAE(nn.Module):

    def __init__(self, latent_size):
```

```

super(ConvVAE, self).__init__()

self.latent_size = latent_size
#encoder
self.encoder = nn.Sequential(
    OrderedDict([
        ('conv1',nn.Conv2d(1, 64, kernel_size=4, stride=2, padding=1)),
        ('relu1',nn.ReLU()),
        ('conv2',nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1)),
        ('relu2',nn.ReLU()),
        ('flatten',Flatten()), #flatten
        ('fc1',nn.Linear(6272, 1024)),
        ('relu3',nn.ReLU())
    ]))
#latent space
# hidden => mu
self.fc1 = nn.Linear(1024, self.latent_size)

# hidden => logvar
self.fc2 = nn.Linear(1024, self.latent_size)
#decoder
self.decoder = nn.Sequential(
    OrderedDict([
        ('fc1',nn.Linear(self.latent_size, 1024)),
        ('relu1',nn.ReLU()),
        ('fc2',nn.Linear(1024, 6272)),
        ('relu2',nn.ReLU()),
        ('unflatten',Unflatten(128, 7, 7)),
        ('relu3',nn.ReLU()),
        ('convT1',nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2,
→padding=1)),
        ('relu4',nn.ReLU()),
        ('convT2',nn.ConvTranspose2d(64, 1, kernel_size=4, stride=2,
→padding=1)),
        ('sigmoid',nn.Sigmoid())
    ]))

def reparameterize(self, mu, logvar):
    if self.training:
        std = torch.exp(0.5 * logvar)
        eps = torch.randn_like(std)
        return eps.mul(std).add_(mu)
    else:
        return mu

def forward(self, x):
    h = self.encoder(x)

```

```

mu, logvar = self.fc1(h), self.fc2(h)
z = self.reparameterize(mu, logvar)
decoded = self.decoder(z)
return decoded, mu, logvar

```

```

[31]: model = ConvVAE(latent_size=2).cuda()
optimizer = optim.Adam(model.parameters(), lr=1e-3)

```

```

[32]: def kld(mu, sigma):
    #return 0.01*0.5*torch.sum(mu**2 + sigma**2-1-torch.log(sigma**2))
    beta=3
    return beta*0.5*torch.sum(mu**2 + sigma**2-1-torch.log(sigma**2))

```

```

[33]: def train(model: nn.Module, optimizer: optim.Optimizer, data: train_loader,
    ↪max_epochs: int, cuda=True):
    total_loss = []
    for epoch in range(max_epochs):
        for batch in data:
            img, _ = batch
            img = Variable(img).cuda()
            # =====forward=====
            reconstructed_img, mu,sigma = model(img)
            kl_loss = kld(mu,sigma)
            reduction_loss = nn.BCELoss(reduction='sum')
            #print(kl_loss, reduction_loss(reconstructed_img,img))
            loss = kl_loss + reduction_loss(reconstructed_img,img)
            # =====backward=====
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            # =====log=====
            total_loss.append(loss.item())
            print('epoch [{}/{}], loss:{:.4f}'.format(epoch+1, max_epochs, loss.
    ↪item()))
    return total_loss

```

```

[34]: VAE_losses = train(model, optimizer, train_loader, max_epochs=30)

```

```


epoch [1/30], loss:-849837.8750
epoch [2/30], loss:-888150.7500
epoch [3/30], loss:-845702.0625
epoch [4/30], loss:-840971.8750
epoch [5/30], loss:-917778.7500
epoch [6/30], loss:-917551.3750
epoch [7/30], loss:-913599.3125
epoch [8/30], loss:-941815.7500
epoch [9/30], loss:-944280.6875

```

```
epoch [10/30], loss:-931594.3750
epoch [11/30], loss:-976067.0000
epoch [12/30], loss:-917829.6250
epoch [13/30], loss:-874786.6250
epoch [14/30], loss:-956616.0000
epoch [15/30], loss:-955056.0000
epoch [16/30], loss:-945033.6875
epoch [17/30], loss:-994633.3750
epoch [18/30], loss:-949189.1250
epoch [19/30], loss:-959227.5625
epoch [20/30], loss:-976263.8750
epoch [21/30], loss:-1008004.5000
epoch [22/30], loss:-1008070.3125
epoch [23/30], loss:-1020752.2500
epoch [24/30], loss:-941680.3125
epoch [25/30], loss:-1022868.2500
epoch [26/30], loss:-955739.0625
epoch [27/30], loss:-971300.9375
epoch [28/30], loss:-1033076.6875
epoch [29/30], loss:-990972.1875
epoch [30/30], loss:-1009383.2500
```

```
[35]: plt.plot(VAE_losses)
```


```
[35]: [<matplotlib.lines.Line2D at 0x7f1cd14cbac8>]
```



```
[36]: test_examples = None

with torch.no_grad():
    for batch_features in test_loader:
        batch_features = batch_features[0]
        test_examples = batch_features.cuda()
        reconstruction, _, _ = model(test_examples)
        break
```


```
[37]: with torch.no_grad():
    number = 10
    plt.figure(figsize=(20, 4))
    plt.suptitle('Original Images', fontsize=16)
    for index in range(number):
        # display original
        ax = plt.subplot(2, number, index + 1)
        plt.imshow(test_examples.cpu()[index].numpy().reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)
    plt.suptitle('Reconstructed Images', fontsize=16)
    for index in range(number):
        # display reconstruction
        ax = plt.subplot(2, number, index + 1 + number)
        plt.imshow(reconstruction.cpu()[index].numpy().reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)
    plt.show()
```



0.2.1 a) Train a Variational Autoencoder with latent dimension of size 2. Then, plot the digits where their associated position was in latent space similarly as explained in the lecture

code taken from here <https://avandekleut.github.io/vae/>

```
[38]: def plot_latent(model, data, num_batches = 100):
    for i, (img, labels) in enumerate(data):
        img = Variable(img).cuda()
        h = model.encoder(img)
        latent_space= model.fc1(h)
        #latent_space, _ = model.encode(img)
        latent_space = latent_space.reshape([64,-1,2])
        latent_space = latent_space.to('cpu').detach().numpy()
        plt.scatter(latent_space[:, :, 0], latent_space[:, :, 1], c=labels.to('cpu').
        →detach().numpy(), cmap='tab10')
        if i > num_batches:
            plt.colorbar()
            break
    plot_latent(model, test_loader)
```



```
[39]: def plot_reconstructed(model, r0=(-5, 10), r1=(-10, 5), n=12):
    w = 28
    img = np.zeros((n*w, n*w))
    for i, y in enumerate(np.linspace(*r1, n)):
        for j, x in enumerate(np.linspace(*r0, n)):
            z = torch.Tensor([[x, y]]).cuda()
            #print(z)
            x_hat = model.decoder(z)
            x_hat = x_hat.reshape(28, 28).to('cpu').detach().numpy()
            img[(n-1-i)*w:(n-1-i+1)*w, j*w:(j+1)*w] = x_hat
```

```
plt.imshow(img, extent=[*r0, *r1])  
plot_reconstructed(model)
```

