

# Logistic\_Regression

November 27, 2020

Group Member Name(Master Data Science): Suresh Kumar Choudhary, Sofya Laskina, Emilio Kuhlmann

Import necessary packages: Numpy, Pandas, matplotlib

```
[59]: import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
```

Mount your google drive (if you have a google account) or upload files (go on the file icon on the left -> right click). Copy path of zip.train and zip.test and load them as numpy arrays using the following code (insert the path as string).

```
[60]: path_to_train = './zip.train'
path_to_test = './zip.test'
training_data = np.array(pd.read_csv(path_to_train, sep=' ', header=None))
test_data = np.array(pd.read_csv(path_to_test, sep=' ', header=None))

X_train, y_train = training_data[:,1:-1], training_data[:,0]
X_test, y_test = test_data[:,1:], test_data[:,0]

# We only want to classify two different digits. You can choose which digits you
→ want to classify yourself

X_train = X_train[np.logical_or(y_train == 1, y_train == 9)]
X_test = X_test[np.logical_or(y_test == 1, y_test == 9)]
y_train = y_train[np.logical_or(y_train == 1, y_train == 9)].reshape(-1,1)
y_train=np.where(y_train> 5, 1, 0)
#print(y_train)
y_test = y_test[np.logical_or(y_test == 1, y_test == 9)].reshape(-1,1)
y_test=np.where(y_test> 5, 1, 0)
```

```
[61]: def show_numbers(X):
    num_samples = 90
    indices = np.random.choice(range(len(X)), num_samples)
    print(indices.shape)
```

```

sample_digits = X[indices]


fig = plt.figure(figsize=(20, 6))

for i in range(num_samples):
    ax = plt.subplot(6, 15, i + 1)
    img = 1-sample_digits[i].reshape((16, 16))
    plt.imshow(img, cmap='gray')
    plt.axis('off')

```

```
[62]: show_numbers(X_test)
```

(90,)



Implement Logistic Regression, do gradient descent until training converges (find a good criterion for when that is the case yourself) and test the accuracy on your test data.

```

[79]: class LogisticRegression:
    def __init__(self, w, b):
        self.w = w
        self.b = b

    def mean_squared_error(self, x, y):
        mse = 1/y.shape[0]*np.sum((y-self.sigmoid_function(x))**2)
        return mse

    def sigmoid_function(self, x):
        return 1 / (1 + np.exp(-(np.dot(x, self.w) + self.b)))

    def cross_entropy_loss(self, x, y):
        sigmoid = self.sigmoid_function(x)
        CE = -1/x.shape[0]*np.sum(np.multiply(y, np.log(sigmoid))+ np.
→multiply((1-y), np.log(1-sigmoid)))
        return CE

```

```

def gradient_descent(self, x,y):
    dw = 1/x.shape[0] * np.dot(x.T, (self.sigmoid_function(x) - y))
    db = 1/x.shape[0] * np.sum(self.sigmoid_function(x) - y)
    return (dw,db)

def fit(self,x,y, learning_rate,num_iter):
    counter = 0
    loss = [self.cross_entropy_loss(x,y)]
    loss_difference = 1
    while loss_difference > 0.001 and counter < num_iter:
        gradient_updated = self.gradient_descent(x,y)
        #print(gradient_updated[0][:10])
        #print(gradient_updated[0].shape[0],self.w.shape[0])
        w_updates = self.w - learning_rate*gradient_updated[0]
        b_updates = self.b - learning_rate*gradient_updated[1]
        self.w = w_updates
        self.b = b_updates
        loss.append(self.cross_entropy_loss(x,y))
        #print(current_loss)
        loss_difference = abs(loss[-2]-loss[-1])

        counter+=1
    self.plot_loss(loss,np.arange(0,counter+1,1),learning_rate)
    return LogisticRegression(self.w,self.b)

#gradient for mean squared error model
def greadient_meansquared(self,x,y):
    dw = 1/x.shape[0] *np.dot(x.T,self.sigmoid_function(x)*(1-self.
→sigmoid_function(x).T)*(self.sigmoid_function(x) - y))
    db = 1/x.shape[0] * np.sum(self.sigmoid_function(x) - y)
    return (dw,db)

#means quared error model fit function
def fit_mse(self,x,y, learning_rate,num_iter):
    counter = 0
    error = [self.mean_squared_error(x,y)]
    error_difference = 1
    while error_difference > 0.001 and counter < num_iter:
        gradient_updated = self.greadient_meansquared(x,y)
        #print(gradient_updated[0][:10])\
        print('gradient shape',gradient_updated[0].shape[0], 'weight shape',self.w.
→shape[0])
        w_updates = self.w - learning_rate*gradient_updated[0]
        b_updates = self.b - learning_rate*gradient_updated[1]
        self.w = w_updates
        self.b = b_updates

```

```

        error.append(self.mean_squared_error(x,y))
        #print(current_loss)
        error_difference = abs(error[-2]-error[-1])

        counter+=1
        self.plot_mse(error,np.arange(0,counter+1,1),learning_rate)
        return LogisticRegression(self.w,self.b)

def predict(self,x):
    y = self.sigmoid_function(x)
    return y
def plot_loss(self,loss,num_iter,learning_rate):
    #print(loss)
    #print(num_iter)
    plt.plot(num_iter,loss)
    plt.title('Learning rate=%f'%learning_rate)
    plt.xlabel('Number of iterations')
    plt.ylabel('loss')
    plt.show()
def plot_mse(self,loss,num_iter,learning_rate):
    #print(loss)
    #print(num_iter)
    plt.plot(num_iter,loss)
    plt.title('Learning rate=%f, Mean Squared Error vs Iterations_
→ '%learning_rate)
    plt.xlabel('Number of iterations')
    plt.ylabel('Mean Squared Error')
    plt.show()

```

```

[64]: def weight_initiliazitation(x):
        w = np.zeros((x.shape[1], 1))
        b = 0
        return (w,b)

def find_optimum_lr(lr_list,X_train,y_train,X_test,y_test,w,b):
    lr_acc=[]
    for lr in lr_list:
        model = LogisticRegression(w,b)
        model.fit(X_train,y_train,lr,num_iter)
        predictions = model.predict(X_test)
        y_pred_binary = np.where(predictions > 0.5, 1, 0)
        #we are using f-score as no of digits are not balanced
        accuracy_score=f1_score(y_test,y_pred_binary)
        lr_acc.append([accuracy_score,lr])
    x=[str(x[1]) for x in lr_acc]
    y=[x[0] for x in lr_acc]
    plt.plot(x,y)

```

```


plt.xlabel('learning rate')
plt.ylabel('Accuracy score(F-score)')
plt.show()
lr_acc.sort(reverse=True)
return lr_acc[0][1]



```



```



[65]: num_iter=500
w,b=weight_initiliazitation(X_train)
lr_list=[0.00001,0.0001,0.001,0.01,0.1,0.2,0.3,0.4]
optimum_learning_rate=find_optimum_lr(lr_list,X_train,y_train,X_test,y_test,w,b)
model = LogisticRegression(w,b).
→fit(X_train,y_train,optimum_learning_rate,num_iter)

```



















```
[66]: model = LogisticRegression(w,b).  
      ↪ fit(X_train,y_train,optimum_learning_rate,num_iter)
```



```
[67]: predictions = model.predict(X_test)
y_pred_binary = np.where(predictions > 0.5, 1, 0)
accuracy_score=f1_score(y_test,y_pred_binary)
print('accuracy score(F-score)%.2f'%accuracy_score)
```


accuracy score(F-score)0.98

```
[68]: plt.hist(predictions, bins = 10)
plt.title('Model Probabilities Distribution')
plt.xlabel('model probabilities')
plt.ylabel('freq')
plt.show()

# Generate a classification report
cm_plot_labels = ['digit 1', 'digit 9']
# For this to work we need y_pred as binary labels not as probabilities
y_pred_binary = np.where(predictions > 0.5, 1, 0)


report = classification_report(y_test, y_pred_binary,
    →target_names=cm_plot_labels)

print(report)
```



	precision	recall	f1-score	support
digit 1	0.98	0.98	0.98	264
digit 9	0.98	0.98	0.98	177
accuracy			0.98	441
macro avg	0.98	0.98	0.98	441
weighted avg	0.98	0.98	0.98	441

```
[69]: from sklearn import metrics
fpr, tpr, _ = metrics.roc_curve(y_test, predictions)
auc = metrics.roc_auc_score(y_test, predictions)
plt.plot(fpr,tpr,label="logistic regression, auc=%.3f"%auc)
plt.legend(loc=4)
plt.title('ROC curve')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.show()
```



Logistic Regression can be interpreted as a neural network with just a single layer. It uses the Cross Entropy to measure the performance of the layer (i.e. of the “trained” weight  $\mathbf{w}$ ). In ML we call this the **Loss function**.

What happens when you take the Means Squared Error (MSE) instead of the Cross Entropy? Does this also work? Implement MSE and try for yourself.

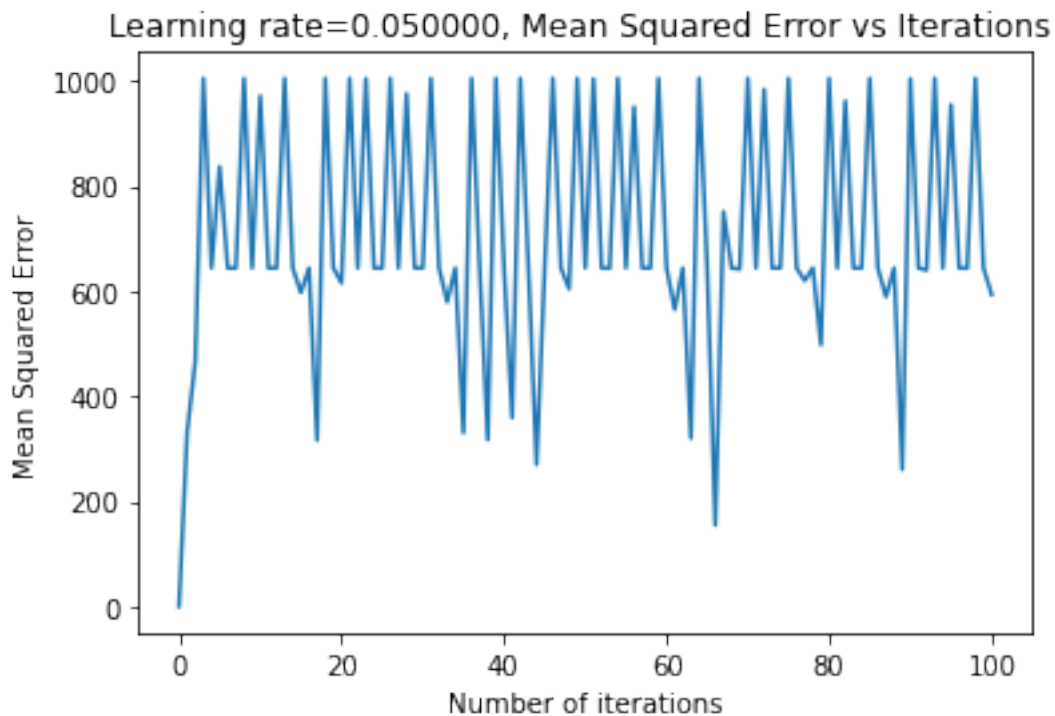
Answer: If we take MSE instead of the Cross Entropy then It will not converge as the cost function is not convex and there are a lot of local minima.

```
[81]: num_iter=100
w,b=weight_initiliazitation(X_train)
#lr_list=[0.00001,0.0001,0.001,0.01,0.1,0.2,0.3,0.4]
#optimum_learning_rate=find_optimum_lr_mse(lr_list,X_train,y_train,X_test,y_test,w,b)
model = LogisticRegression(w,b).fit_mse(X_train,y_train,0.05,num_iter)
```

[illegible]

[illegible]

```
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
gradient shape 256 weight shape 256
```



[ ]: